

```
In [1]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import SGDRegressor
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_squared_error,r2_score,mean_absolute_error
import matplotlib.pyplot as plt
```

```
In [2]: data=pd.read_csv('CarPrice_Assignment.csv')
print(data.head())
print(data.info())
```

```

car_ID symboling                                CarName fueltype aspiration doornumber
 \
0      1          3      alfa-romero giulia      gas      std      two
1      2          3      alfa-romero stelvio     gas      std      two
2      3          1  alfa-romero Quadrifoglio    gas      std      two
3      4          2      audi 100 ls            gas      std      four
4      5          2      audi 100ls           gas      std      four

carbody drivewheel enginelocation wheelbase ... enginesize \
0 convertible       rwd        front     88.6   ...     130
1 convertible       rwd        front     88.6   ...     130
2 hatchback         rwd        front     94.5   ...     152
3 sedan             fwd        front     99.8   ...     109
4 sedan             4wd        front     99.4   ...     136

fuelsystem boreratio stroke compressionratio horsepower peakrpm citympg
 \
0 mpfi        3.47    2.68          9.0      111    5000     21
1 mpfi        3.47    2.68          9.0      111    5000     21
2 mpfi        2.68    3.47          9.0      154    5000     19
3 mpfi        3.19    3.40         10.0      102    5500     24
4 mpfi        3.19    3.40          8.0      115    5500     18

highwaympg price
0      27  13495.0
1      27  16500.0
2      26  16500.0
3      30  13950.0
4      22  17450.0

[5 rows x 26 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
 #   Column          Non-Null Count  Dtype  
--- 
0   car_ID          205 non-null    int64  
1   symboling       205 non-null    int64  
2   CarName         205 non-null    object  
3   fueltype        205 non-null    object  
4   aspiration      205 non-null    object  
5   doornumber      205 non-null    object  
6   carbody         205 non-null    object  
7   drivewheel      205 non-null    object  
8   enginelocation   205 non-null    object  
9   wheelbase       205 non-null    float64
10  carlength       205 non-null    float64
11  carwidth        205 non-null    float64
12  carheight       205 non-null    float64
13  curbweight      205 non-null    int64  
14  enginetype      205 non-null    object  
15  cylindernumber   205 non-null    object  
16  enginesize       205 non-null    int64  
17  fuelsystem       205 non-null    object  
18  boreratio        205 non-null    float64
19  stroke           205 non-null    float64
20  compressionratio  205 non-null    float64
21  horsepower        205 non-null    int64  
22  peakrpm          205 non-null    int64  
23  citympg          205 non-null    int64  
24  highwaympg        205 non-null    int64  
25  price            205 non-null    float64

```

```
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
None
```

In [3]:

```
data=data.drop(['CarName','car_ID'],axis=1)
data = pd.get_dummies(data, drop_first=True)
```

In [4]:

```
X=data.drop('price', axis=1)
Y=data['price']
```

In [5]:

```
scaler = StandardScaler()
X=scaler.fit_transform(X)
Y=scaler.fit_transform(np.array(Y).reshape(-1, 1))
```

In [6]:

```
X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.2,random_state=42)
```

In [7]:

```
sgd_model= SGDRegressor(max_iter=1000, tol=1e-3)
```

In [8]:

```
sgd_model.fit(X_train, Y_train)
```

```
C:\ProgramData\anaconda3\lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected.
Please change the shape of y to (n_samples, ), for example using ravel().
y = column_or_1d(y, warn=True)
```

Out[8]:

```
SGDRegressor
SGDRegressor()
```

In [9]:

```
y_pred = sgd_model.predict(X_test)
```

In [10]:

```
mse = mean_squared_error(Y_test, y_pred)
r2=r2_score(Y_test,y_pred)
mae= mean_absolute_error(Y_test, y_pred)
```

In [11]:

```
print('Name:RAGHUL.S ')
print('Reg no: 212225040325')
print("Mean Squared Error:",mse)
print("Mean Absolute Error:",mae)
print("R-Squared Score:",r2)
```

```
Name:RAGHUL.S
Reg no: 212225040325
Mean Squared Error: 0.20941302034292564
Mean Absolute Error: 0.3137138365510395
R-Squared Score: 0.8315273017580898
```

```
In [12]: print("\nModel Coefficients:")
print("Coefficients:",sgd_model.coef_)
print("Intercept:",sgd_model.intercept_)
```

```
Model Coefficients:
Coefficients: [ 0.03686206  0.10319047  0.00533106  0.17993235  0.01682659  0.1
7919498
 0.29841084 -0.02525896 -0.08568881 -0.00261956  0.10372634  0.03949923
-0.03493016 -0.05170662 -0.01670314  0.02001822  0.01542087 -0.02299429
-0.09335112 -0.01847532 -0.09718638 -0.03063997  0.06334887  0.21566676
 0.01382336 -0.10354977  0.07434479 -0.01558528 -0.0034466  0.00955538
-0.04232329 -0.1525486 -0.08306056  0.00068997 -0.02962348  0.00955538
-0.0055237 -0.01610782  0.01670314 -0.01965514 -0.03575729 -0.04362516
-0.01766087]
Intercept: [-0.00986063]
```

```
In [13]: plt.scatter(Y_test,y_pred)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual vs Predicted Prices using SGD Regressor")
plt.plot([min(Y_test),max(Y_test)],[min(Y_test),max(Y_test)],color='red')
plt.show()
```

